BUILDING A SERVERLESS URL SHORTENER ON AWS

Imagine a Giant Library full of books with very long titles. But reading of long titles is tiring so the Librarian gives each book a short code instead. Now, if you want a book, you just tell the Librarian the short code and they find the full title for you. That is how a severless URL Shortener works.
A serverless URL is a simple tool that takes long URL and creates a short version. When one clicks the short link, they are redirected to the original long URL. A serverless system means you don't need to manage servers. Instead, AWS handles everything for you and you only pay for what you use. The URL shortener consists of:
Frontend (S3 + CloudFront): A simple UI for users to input long URLs and get short links.
API (API Gateway + Lambda): Handles URL shortening and redirection.
Database (DynamoDB): Stores mappings between short and long URLs.
CloudFront: A Content Delivery Network (CDN) that helps deliver contents like websites, images, APIs etc faster to users .

Here, the DynamoDB is like the **Library Catalog** where we store short codes and full book titles (URLs). The **Librarian** is AWS Lambda, which is the brain behind everything. When someone brings a new book (long URL), the Librarian writes a short code in the Catalog and when someone asks for a book by its short code, the Librarian looks up the full title and gives it to them.
The API Gateway acts as the **Request Counter**. The Librarian needs a way to receive a request. API Gateway is like the **Help Desk** in the Library where people ask for books.
S3 acts as the Bookshelf where special books are stored. Some books are really popular, so instead of looking them up every time, the Librarian puts few copies on the shelf (S3) and if someone asks for one, they can get it instantly.
CloudFront acts as a **Messenger**. If too many people keep asking for the same book, the Librarian gets tired. So, the Librarian hires a Messenger (CloudFront). If one asks for a book, the Messenger might already have a copy and deliver it fast.

This guide walks through the process of building a URL shortener using AWS services, including AWS Lambda, API Gateway, DynamoDB, S3 and Cloudfront.

PREREQUISITES

AWS Account

IAM Role with permissions for Lambda, DynamoDB, API Gateway, S3 and CloudFront

STEP 1: CREATE AN IAM ROLE FOR LAMBDA

Since AWS Lambda will interact with multiple services, we need to create an IAM role with the necessary policies.

- Go to IAM in the AWS Console
- Click on Roles
- Create Role
- Select AWS Service
- Choose Lambda
- Click Next
- Attach the necessary policies
- AmazonDynamoDBFullAccess : for reading and writing URLs in DynamoDB
- AWSLambdaBasicExecutionRole: for Lambda logging to CloudWatch
- AmazonS3FullAccess : for reading and writing from S3
- CloudFrontFullAccess : for managing CloudFront distribution
- Click **Next**
- Add a name "**admin-lambda-role**"
- Click **Create Role.**

STEP 2: SETTING UP THE BACKEND (Lambda and DynamoDB)

We need a database to store both short and long URLs

- **Create a DynamoDB Table**
- Go to AWS Console
- Open AWS DynamoDB
- Click create table
- Set up Table Details
  - Table Name: studentData
  - Partition Key: studentid
  - Leave it as string

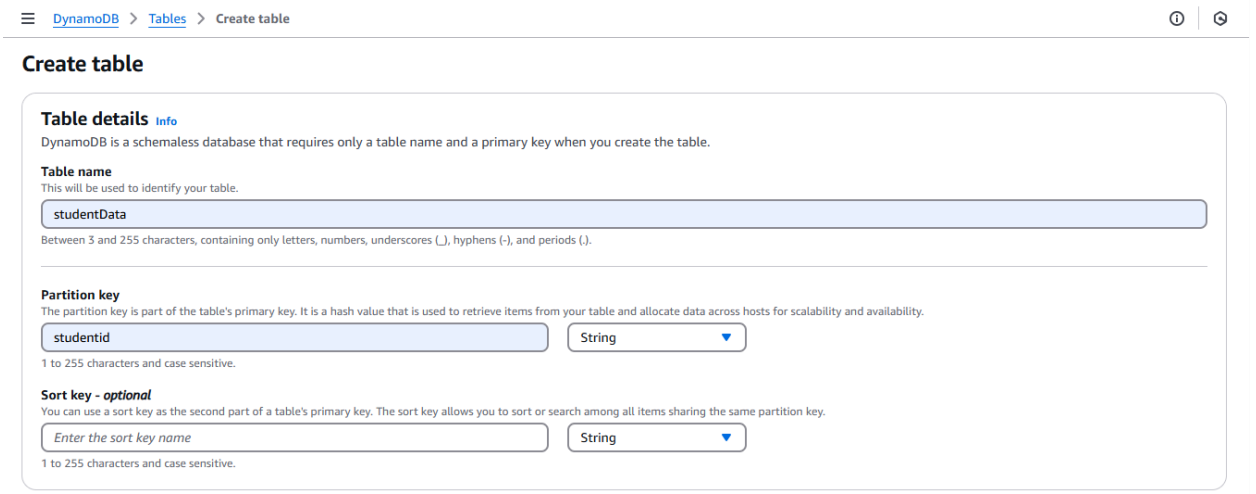

- Leave other settings as default
- Create Table

Now you have a database to store URLs.

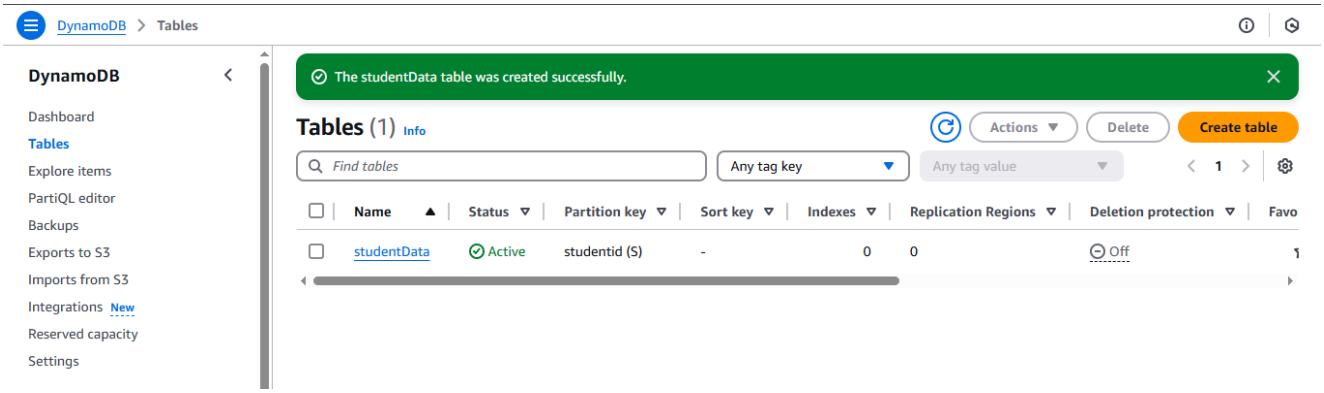

- **Create a Lambda Function**

Lambda will handle URL shortening and redirection.

- Go to AWS Console
- Open AWS Lambda
- Click **Create function**
- Select **Author from scratch**
- Configure Lambda
  - Function Name: **getStudentData**
  - Runtime: **python 3.13**

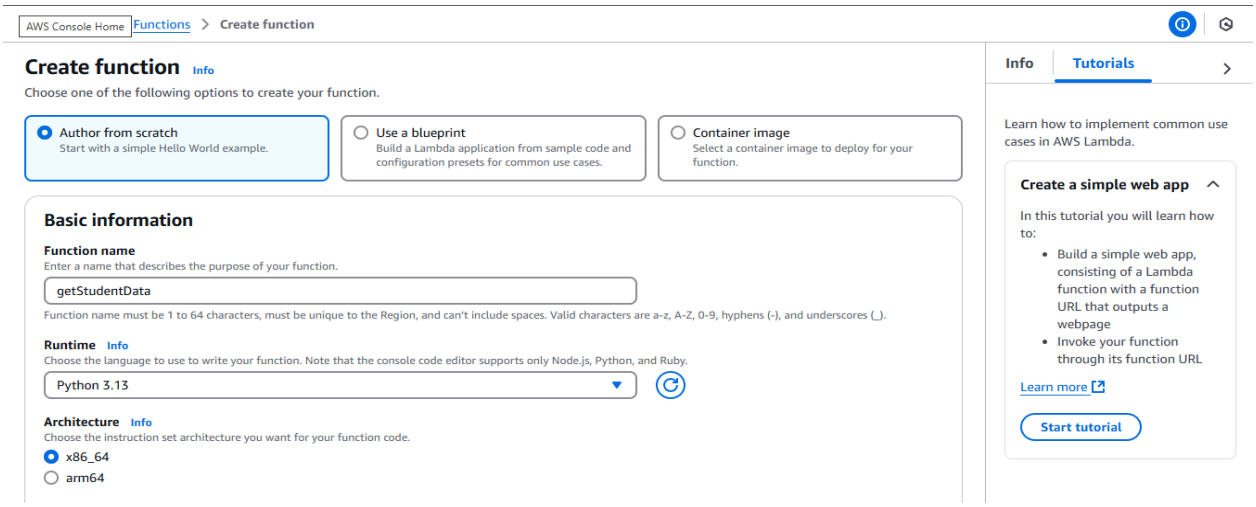

- Permissions: choose use an existing role as a default execution role(select the Lambda role you had created)
- Click **Create Function**.

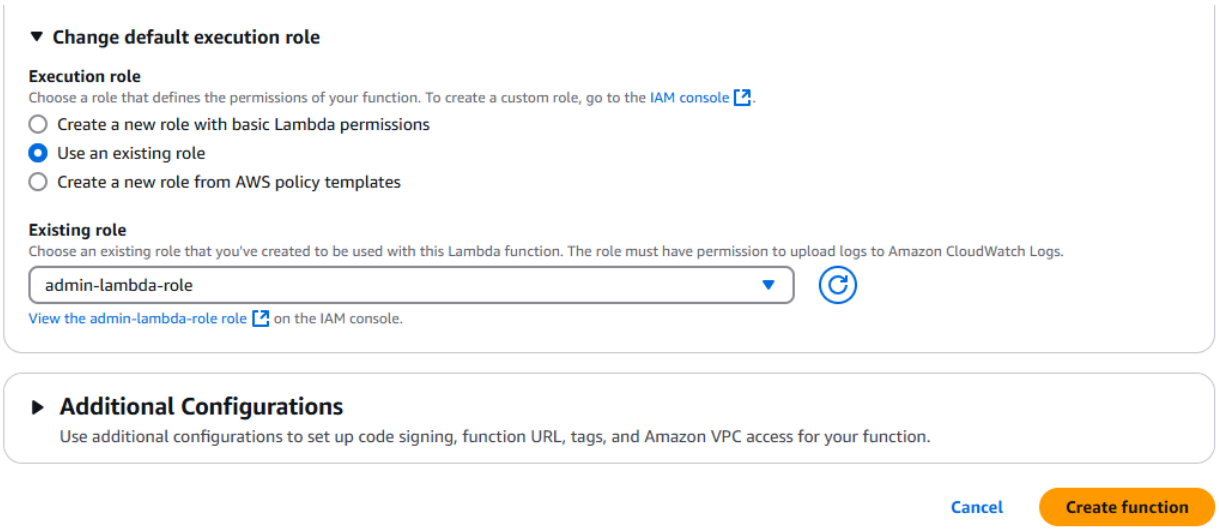

Our Lambda function is ready.

- Add python code for Lambda
- Once the function is created
- Scroll down to **code source**
- Delete the default code and paste the getStudentsData.py code

```
import json
import boto3

def lambda_handler(event, context):
    # Initialize a DynamoDB resource object for the specified region
    dynamodb = boto3.resource('dynamodb', region_name='us-east-2')
```

```
# Select the DynamoDB table named 'studentData'
table = dynamodb.Table('studentData')

# Scan the table to retrieve all items
response = table.scan()
data = response['Items']

# If there are more items to scan, continue scanning until all items are retrieved
while 'LastEvaluatedKey' in response:
    response = table.scan(ExclusiveStartKey=response['LastEvaluatedKey'])
    data.extend(response['Items'])

# Return the retrieved data
return data
```

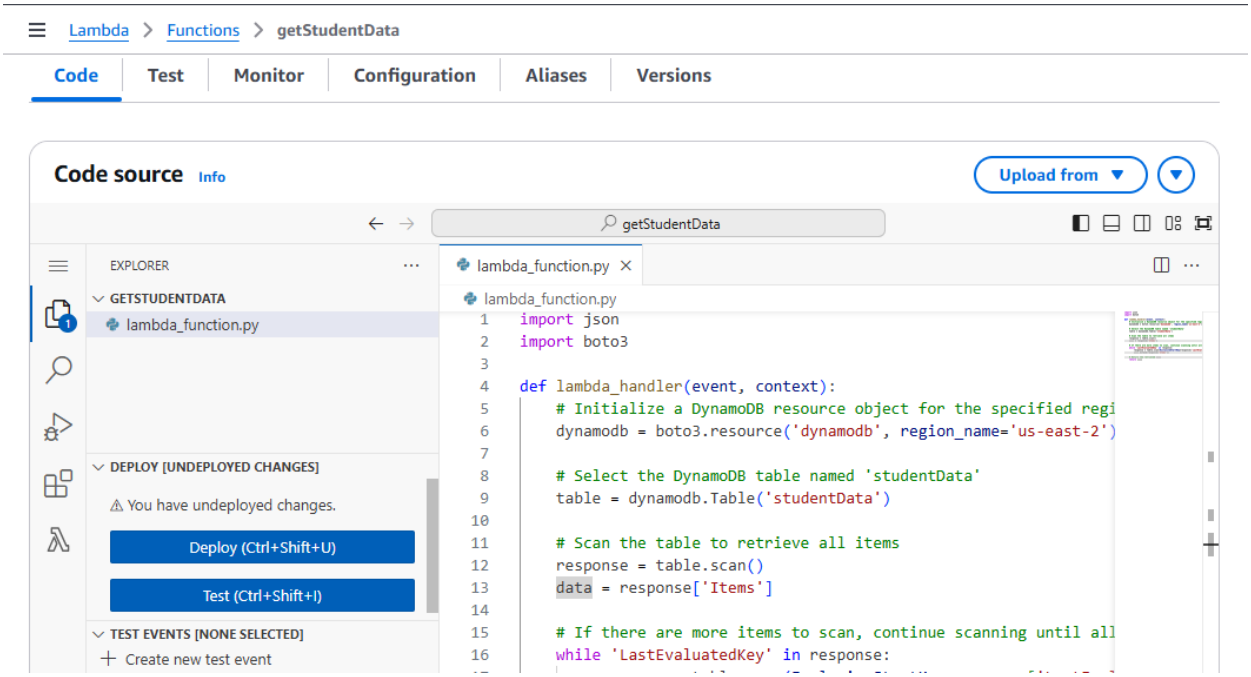

- Deploy Lambda by Clicking **Deploy**

Your Lambda function is now deployed

If you check DynamoDB, you will realize that there is no table.

- Go to Lambda function
- Create a new test event with the name "**mytest**".
- Paste the code on the Event JSON

{
"Key 1": " Value1",
"Key 2": "value2",
"Key 3": " value3"
}

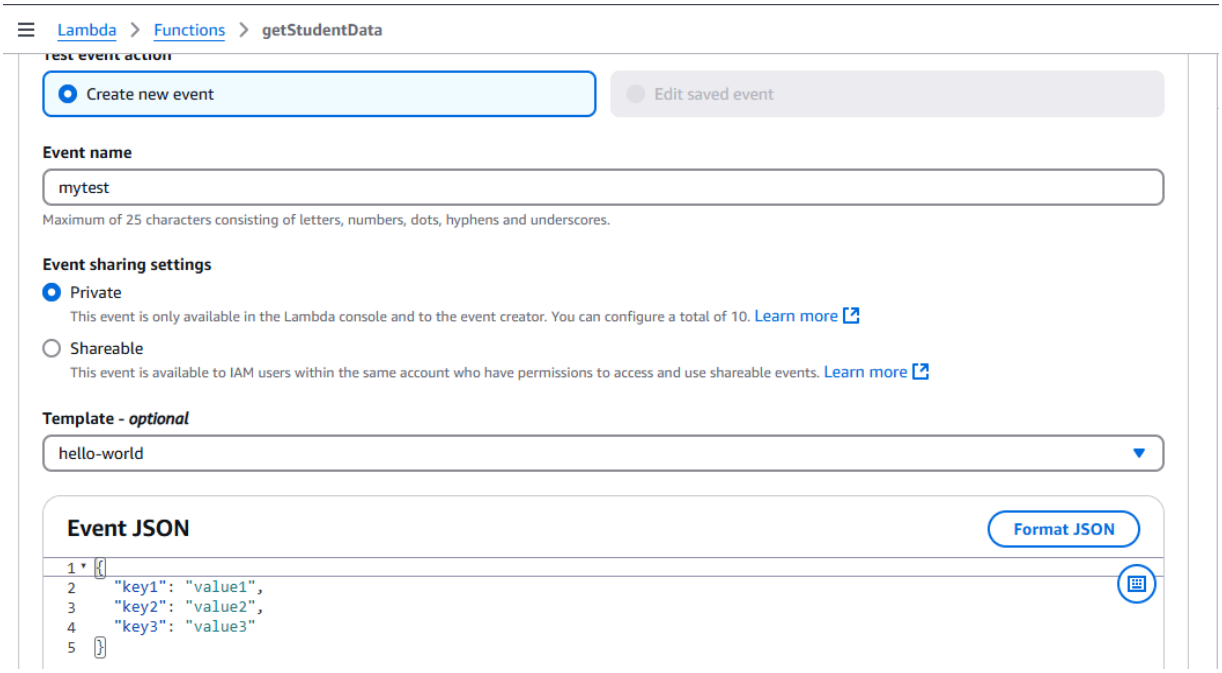

- Save and test the "**newtest**" to invoke the Lambda function.
  This Lambda function will go to DynamoDB Table and try to retrieve data from the table but the response will give empty list so we will create another function for posting data and name it **insertStudentData**.

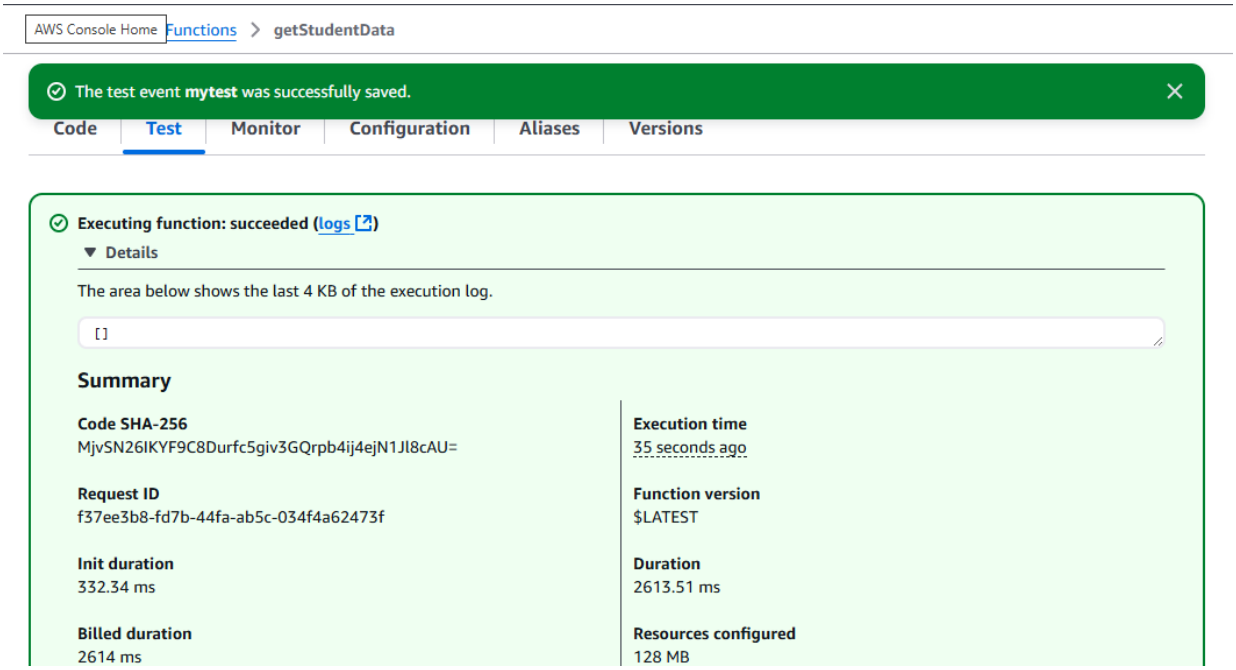

- Go to functions and create a new function with the following details
  - Name: **insertStudentData**
  - Runtime: **python 3.13**

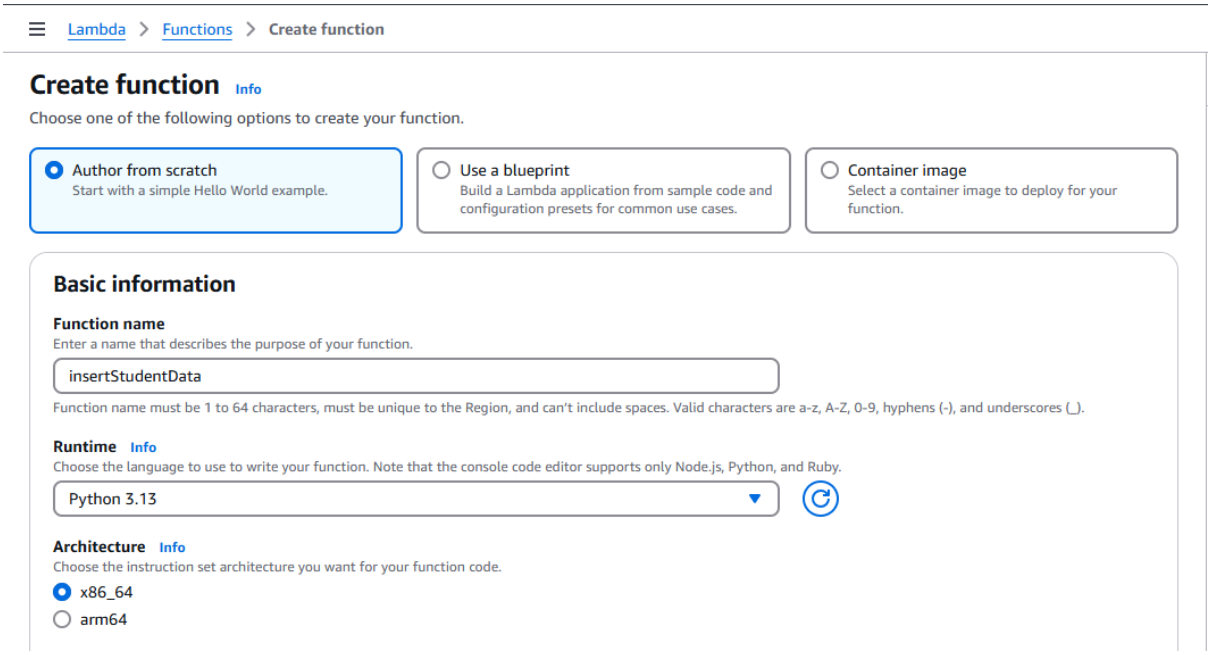

- Permissions: **choose create an existing role** (select the Lambda role you had created)

- Click **Create Function**.
- Copy and paste **insertStudentData.py** file and paste

```
import json
import boto3

# Create a DynamoDB object using the AWS SDK
dynamodb = boto3.resource('dynamodb')
# Use the DynamoDB object to select our table
table = dynamodb.Table('studentData')

# Define the handler function that the Lambda service will use as an entry point
def lambda_handler(event, context):
    # Extract values from the event object we got from the Lambda service and store in
variables
    student_id = event['studentid']
    name = event['name']
    student_class = event['class']
    age = event['age']

    # Write student data to the DynamoDB table and save the response in a variable
    response = table.put_item(
        Item={
            'studentid': student_id,
            'name': name,
            'class': student_class,
            'age': age
        }
    )

    # Return a properly formatted JSON object
    return {
        'statusCode': 200,
        'body': json.dumps('Student data saved successfully!')
    }
```

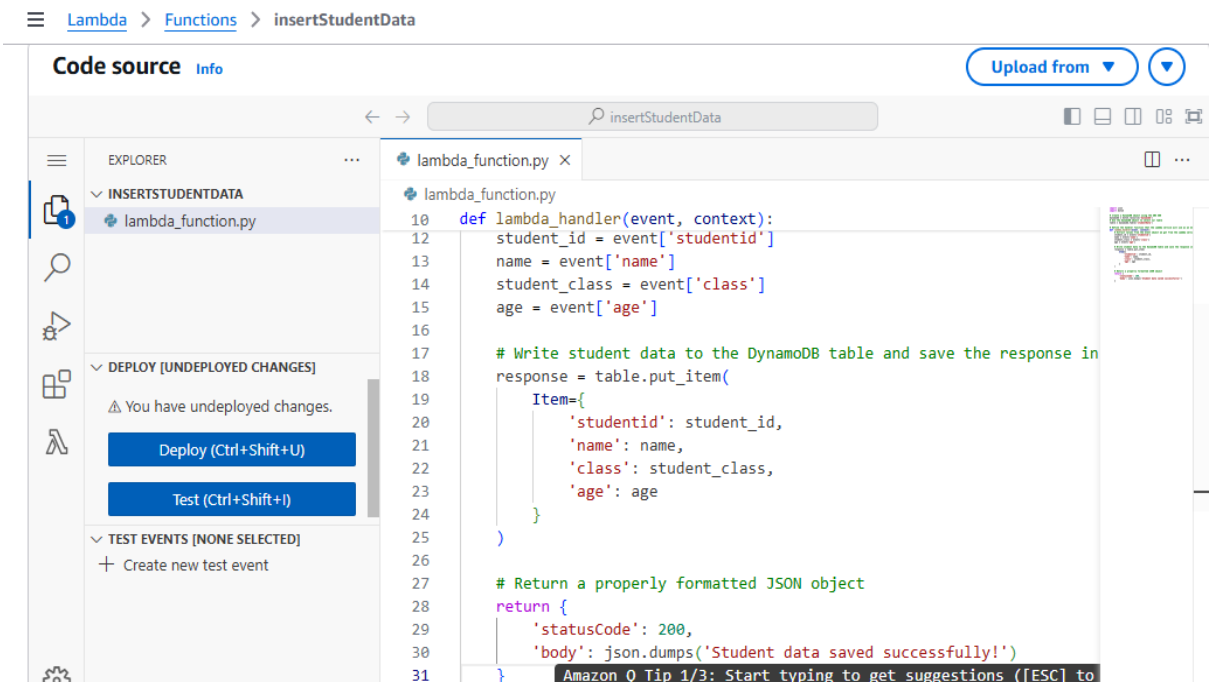

- Click on **deploy**
- Create test event with the name "**mytest**"
  We need to give students name, class, age and id
- Paste the following values. These are the data we will pass to Lambda function

  {
  "studentid": "1",
  "name": "Ekom",
  "class": "A",
  "age": "22"
  }

  This is a part gotten from the inserStudentData.py that has been deployed

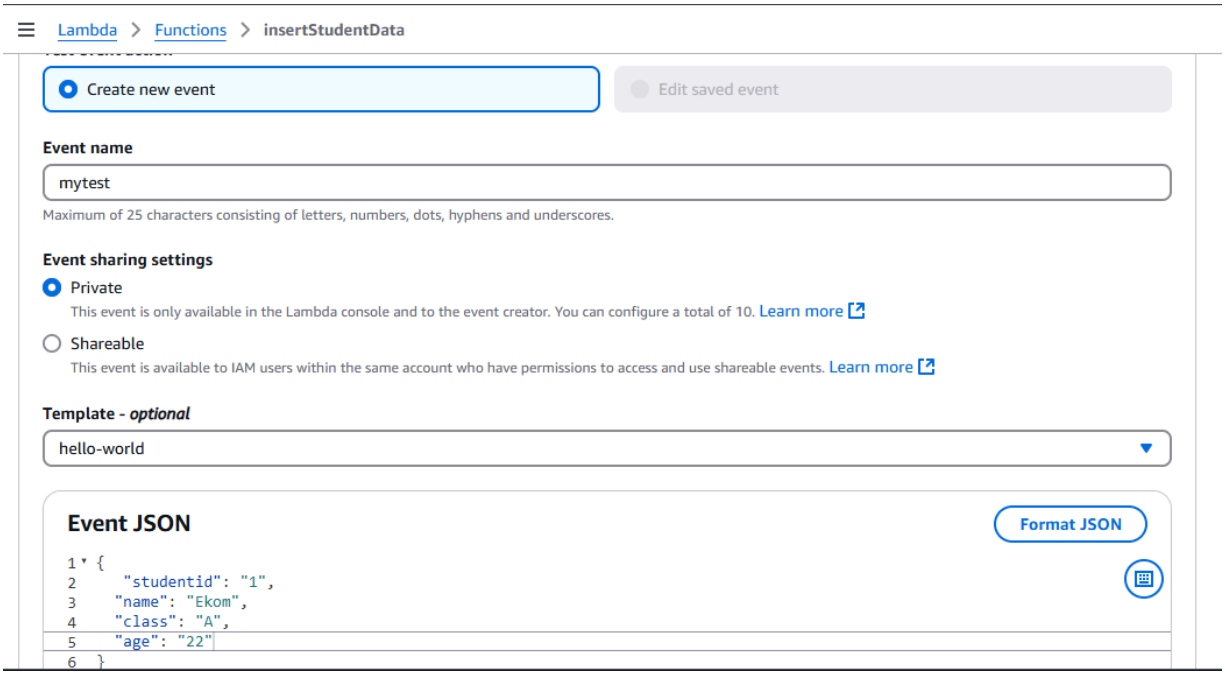

- Save and test

You will see a response "**student data saved successfully**!".

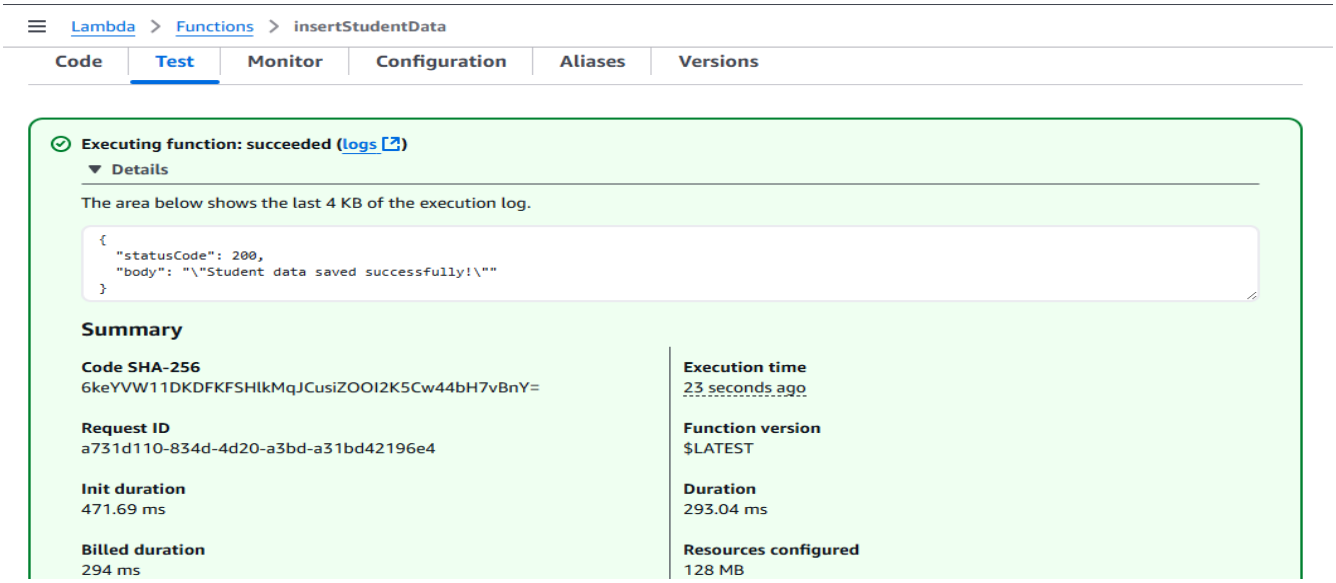

If you go to DynamoDB table and refresh, you will find the student data you had tested on Lambda.

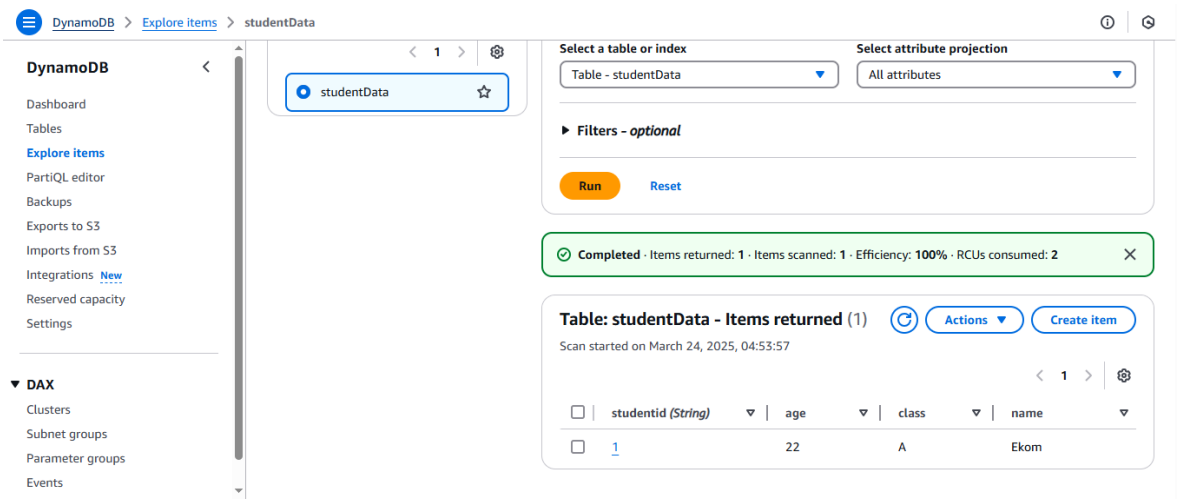

STEP 3: DEPLOY API GATEWAY (FOR PUBLIC ACCESS)

- ➢ Create API Gateway (For public access)

We need an API to connect to Lambda

- Go to API Console
- Open AWS API Gateway
- Click Create API
- Set up API
- Select **Rest API**
- Click Build

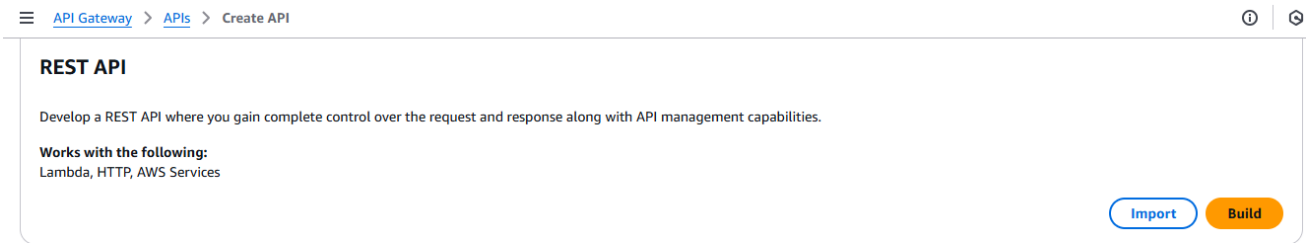

- Input the API details
- Choose New API
- Input API Name as **student**
- Endpoint: Edge-optimized
  Edge-optimized allows users from around the world to access the website anywhere.

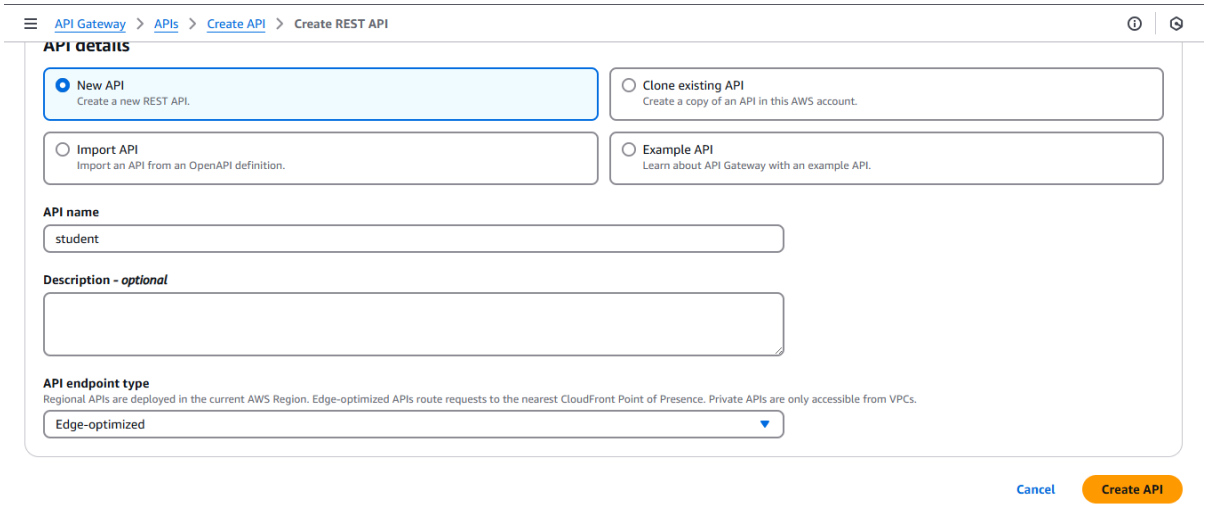

- Create API
- ➢ Create API resources and methods
- Click **create methods**
- Select "**GET**" as type
- Select **Lambda function** as the integration type.
- Select "**getStudentData**" function to connect to your Lambda function.

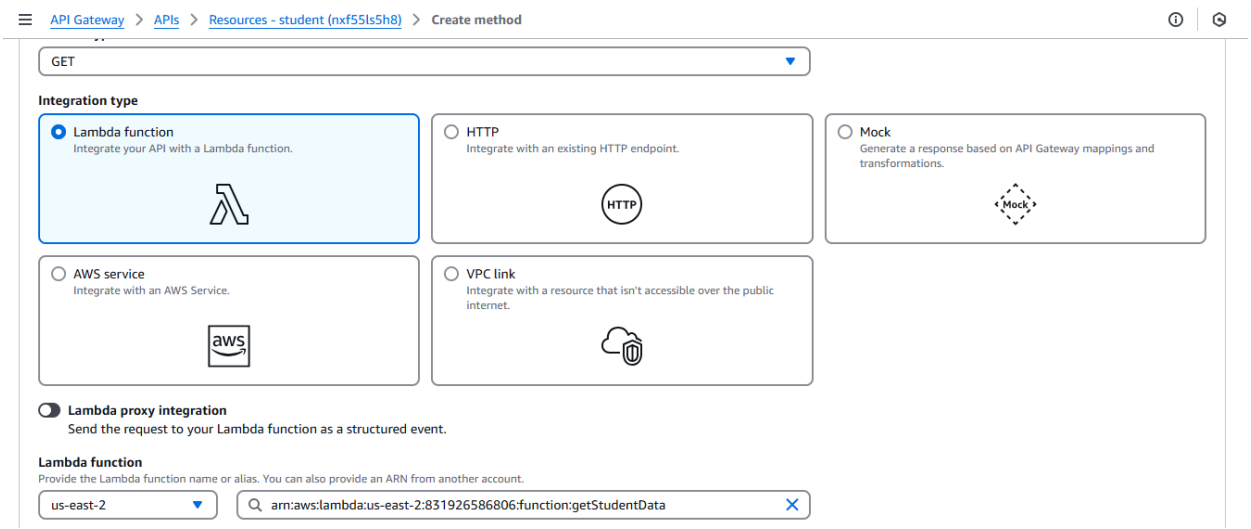

- Click on **Create method**
  We have to create another method
- Select "**POST**" as type
- Select "**insertStudentData**" function to connect to your Lambda function

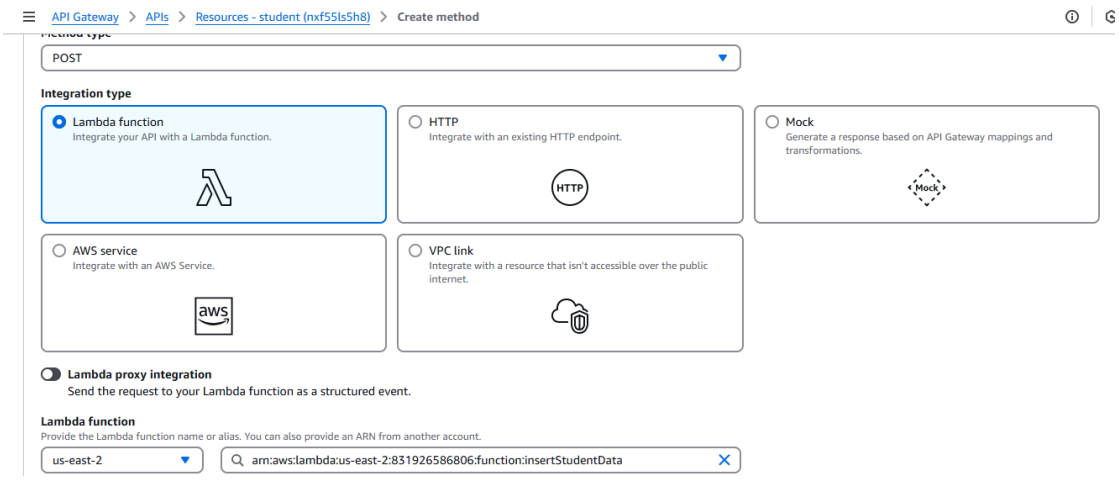

- Click **Deploy API**

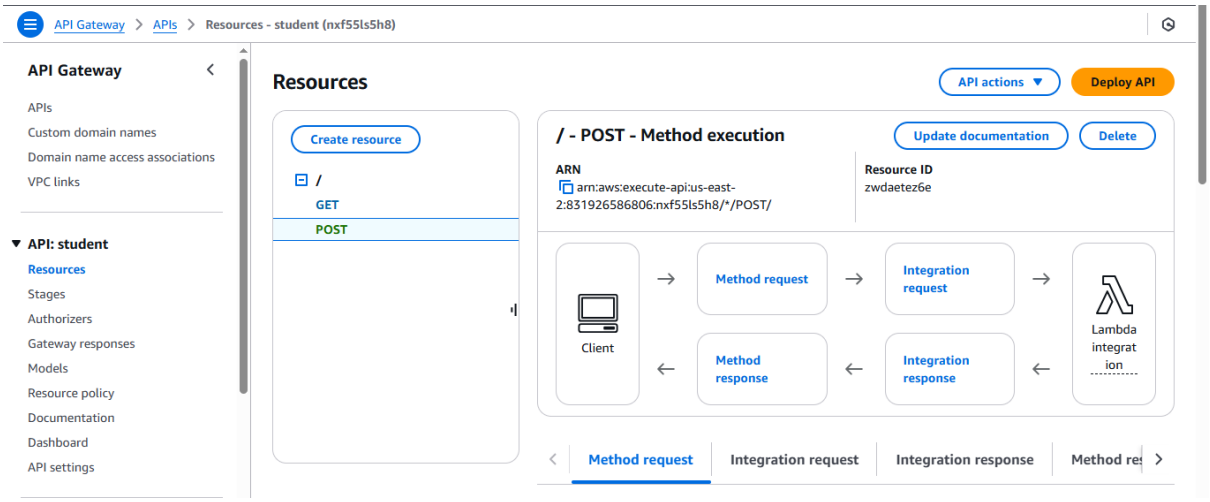

- Select new stage name and type in “**prod**” as the stage name

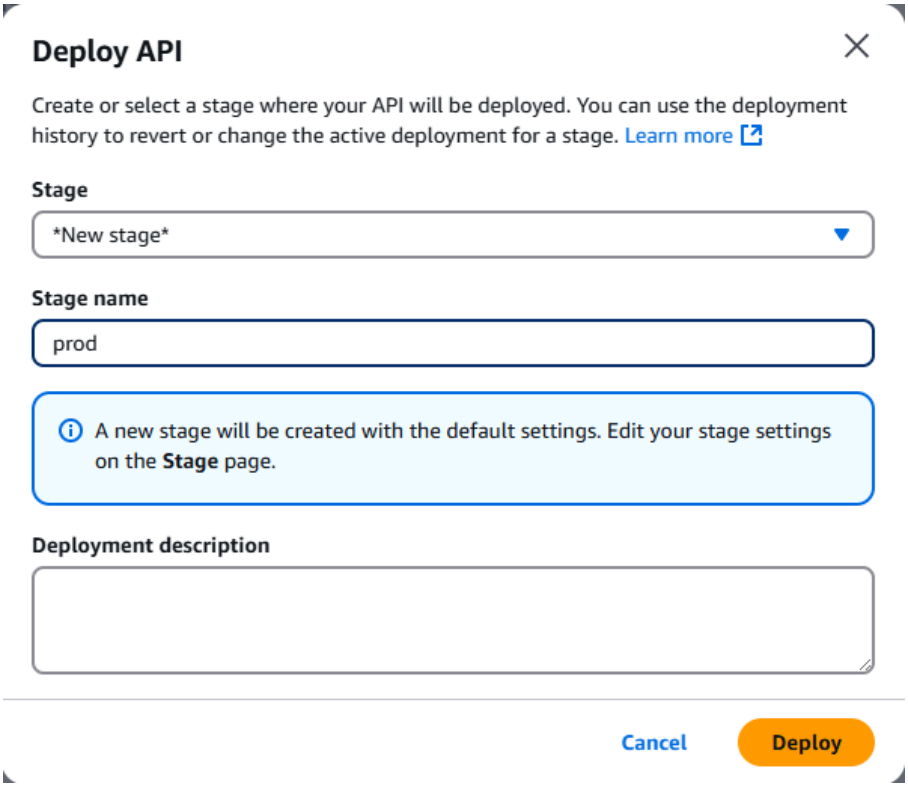

- Click **Deploy**

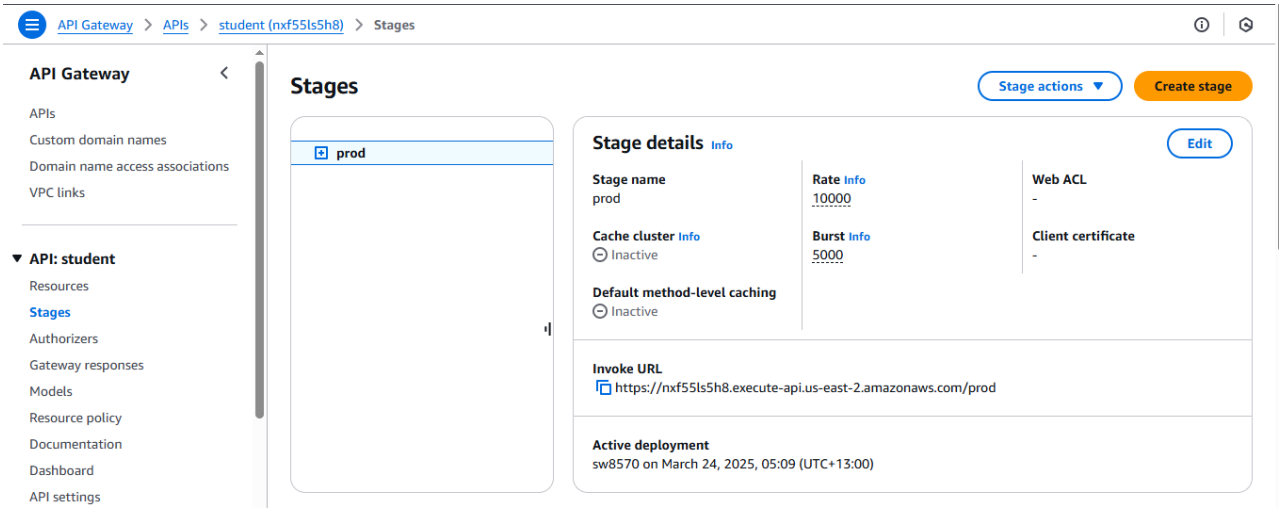

You will find your invoke URL. This URL will invoke our Lambda function.

- Go to resources and enable CORS by enabling **POST** and **GET** in the CORS settings.

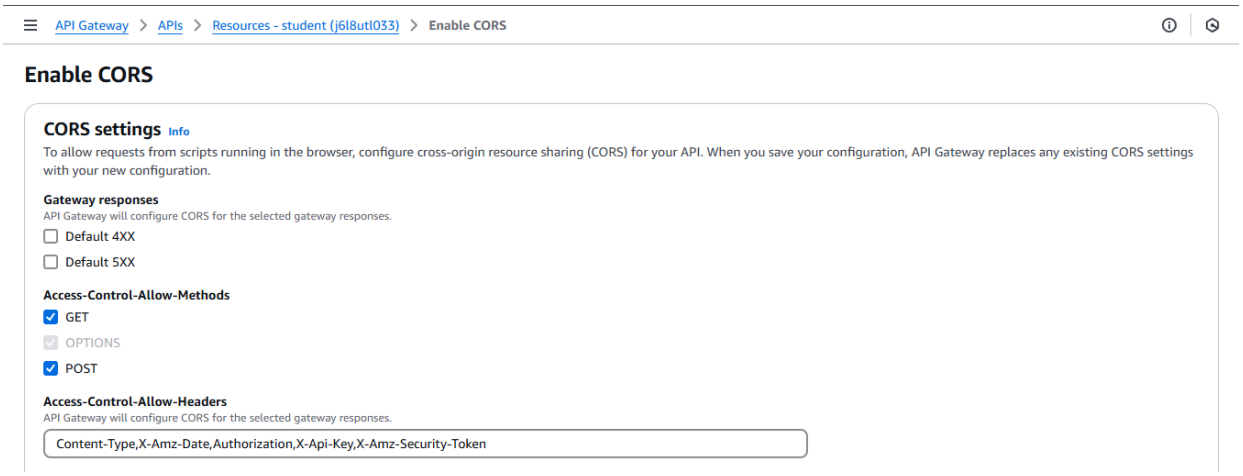

Now your API is live.

- Note the API Gateway URL.
- Copy the invoked URL
- Open the script.js and paste it where you find the endpoint

```
  // Add your API endpoint here
var API_ENDPOINT = "https://j6l8utl033.execute-api.us-east-2.amazonaws.com/prod";
```

When you click on "**getStudentData**", the URL pasted on the .js script will be invoked and it will go to get Lambda function which will be triggered from the DynamoDB table.

STEP 4: DEPLOY THE FRONTEND USING S3 AND CLOUDFRONT

- Go to AWS S3 Console
- Click **Create Bucket**
- Enter a unique Bucket Name **devopppsbucket**

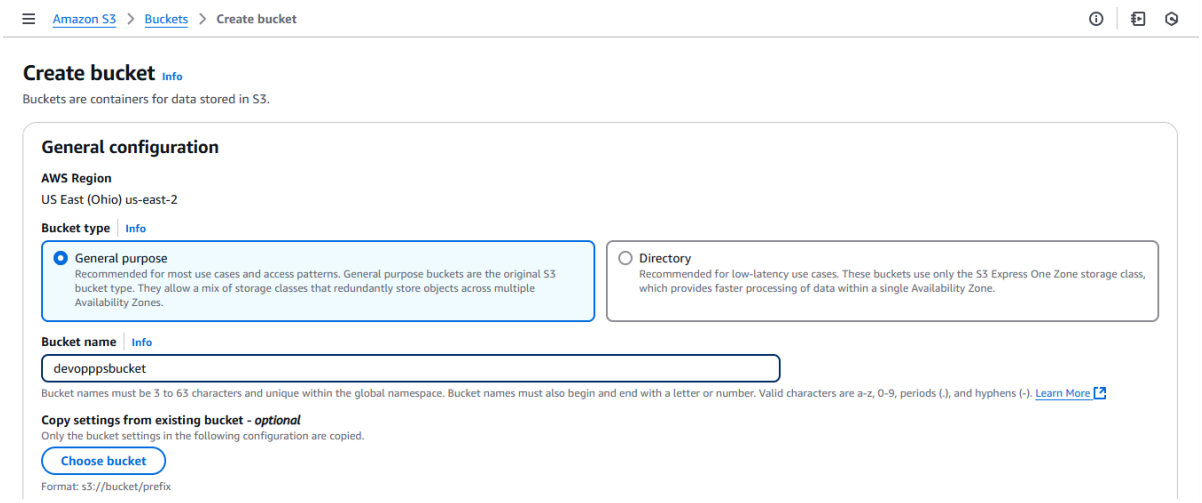

- Uncheck "**Block all public access**" settings
- Click **create bucket** to finalize the setup

➢ Upload Files to the Bucket

- Open the javascript file and replace "Endpoint" with the API Gateway invoke URL
- Click on the Bucket name to open the Bucket.
- Click Upload to select files from your local machine and click **upload.**
- Upload the websites files (an html and javascript frontend for users to enter URLs).

The html is the main web page and the javascript is used to call the API Gateway.

- Upload files to the S3 bucket.

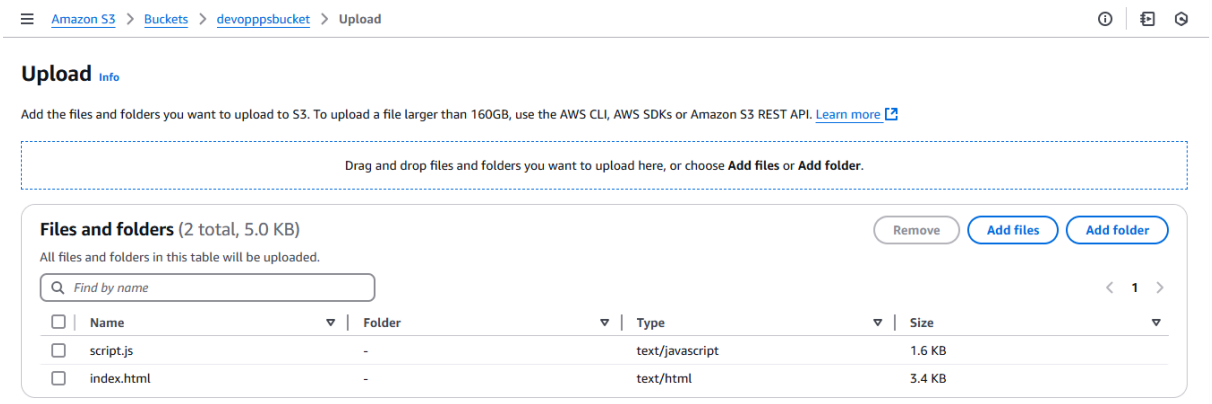

➢ Configure Bucket Permissions (for public access)

- Go to the **Permission** tab of your Bucket.
- Scroll down to the **Bucket Policy** Section and click on Edit to add a policy. Use the Bucket policy.

```
{
   "Version": "2012-10-17",
   "Statement": [
      {
         "Sid": "PublicReadGetObject",
         "Effect": "Allow",
         "Principal": "*",
         "Action": "s3:GetObject",
         "Resource": "arn:aws:s3:::devopppsbucket/*"
      }
   ]
}
```

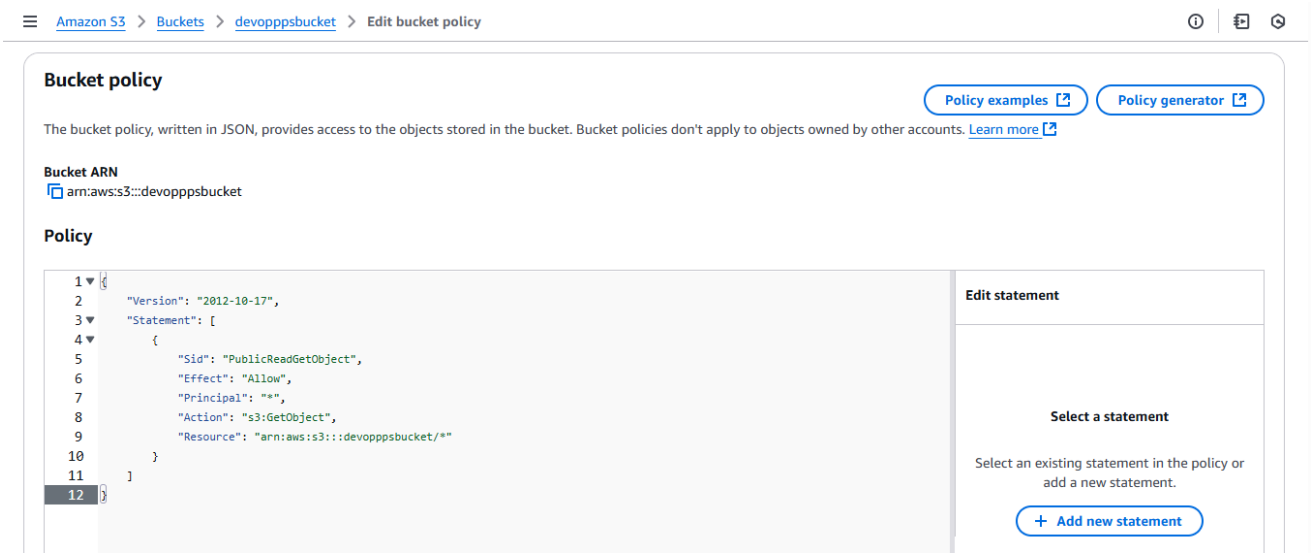

- Click **Save changes** to apply the policy.

➢ Configure Permissions for your Bucket

- Navigate to **Permissions** tab of your Bucket
- Uncheck "**Block all public access**".

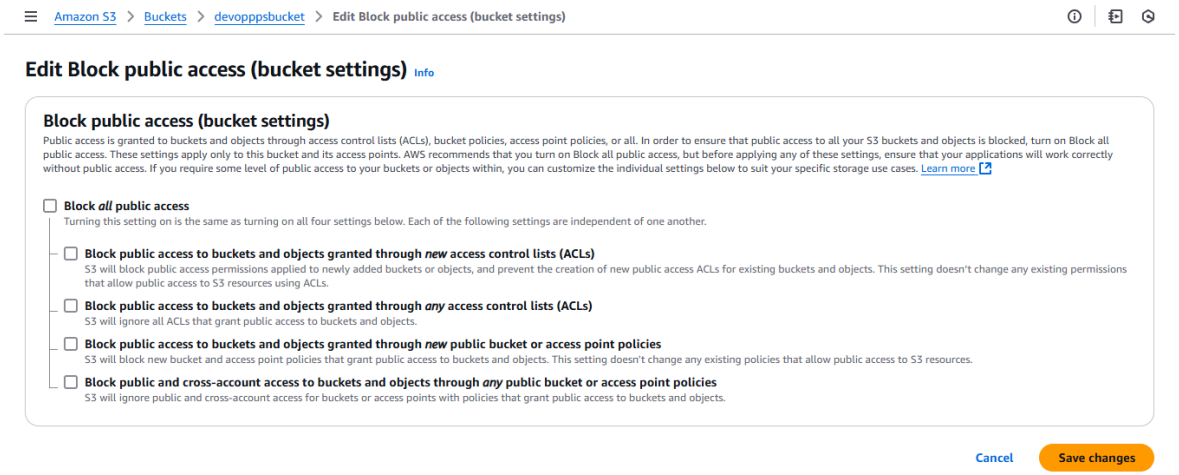

- Click on **Save changes.**

➢ Enable Static Website Hosting

- Go to Properties
- Scroll down to **Static website hosting** and click **Edit**
- Select **Enable**
- Choose **Host a static website.**
- Set index document as **index.html**

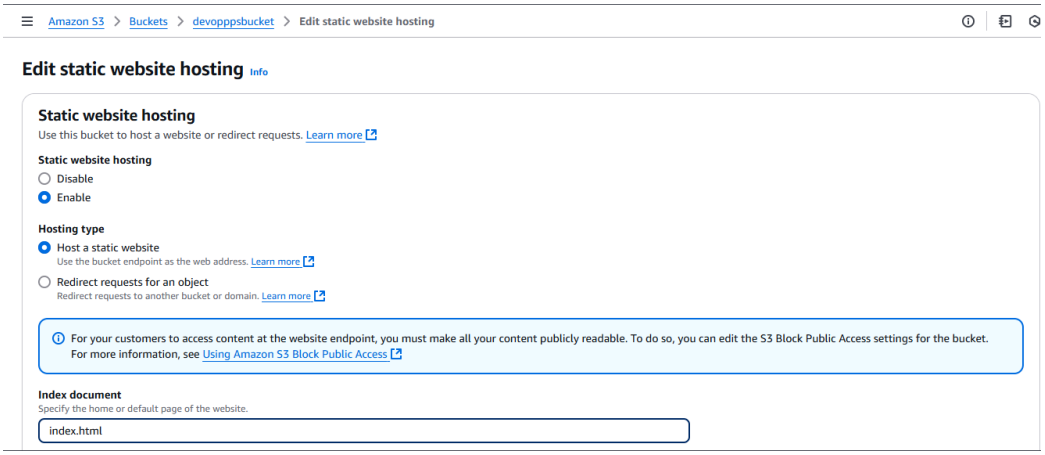

- Click **Save changes**

Your website endpoint will be displayed

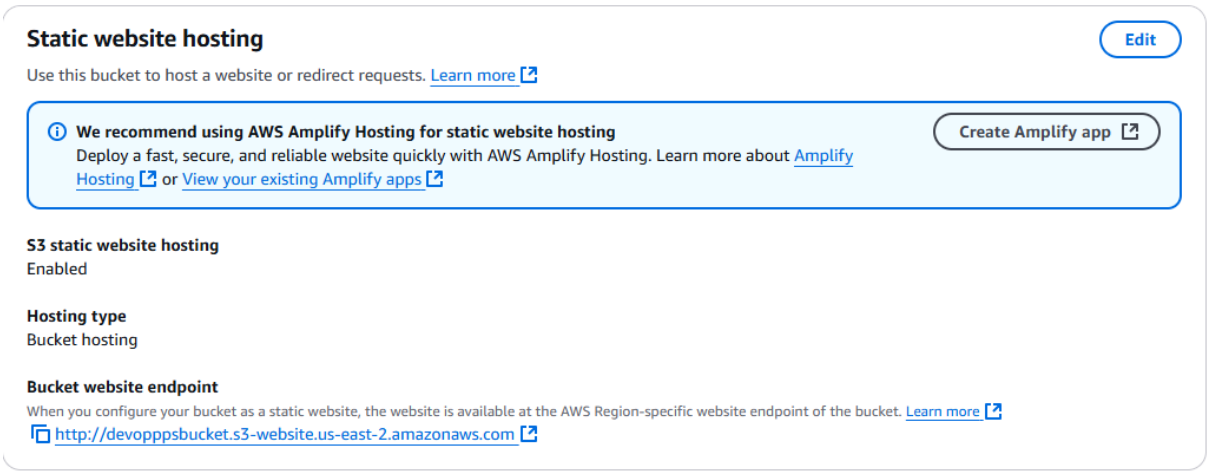

Your application is deployed.

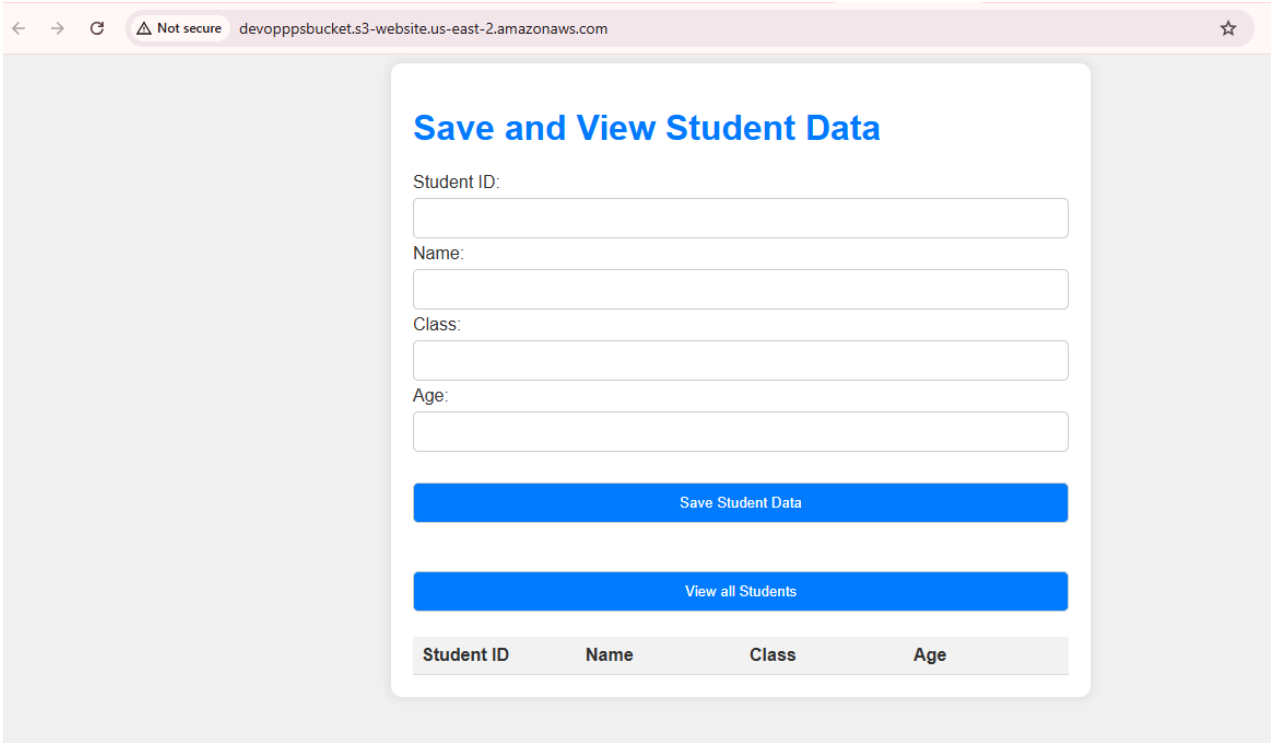

CloudFront would be placed in front of S3 bucket because if you check the website, it is not secure. Its only http which makes our bucket exposed to the public. We need CloudFront to make it secure.

STEP 5: SET UP CLOUDFRONT FOR SHORT URLS

CloudFront will serve as a CDN layer for the URL shortener therefore improving performance.To make the short URL load faster globally and secured, use a custom domain.

- Create a Distribution
  - Go to **CloudFront**
  - Click **Create Distribution**.
  - Select your S3 Bucket as the origin domain

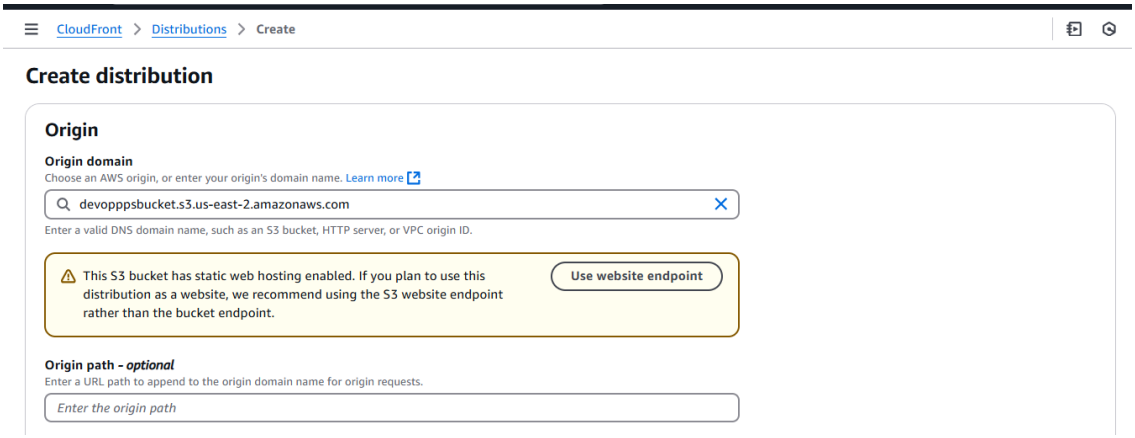

  - Select origin access control.

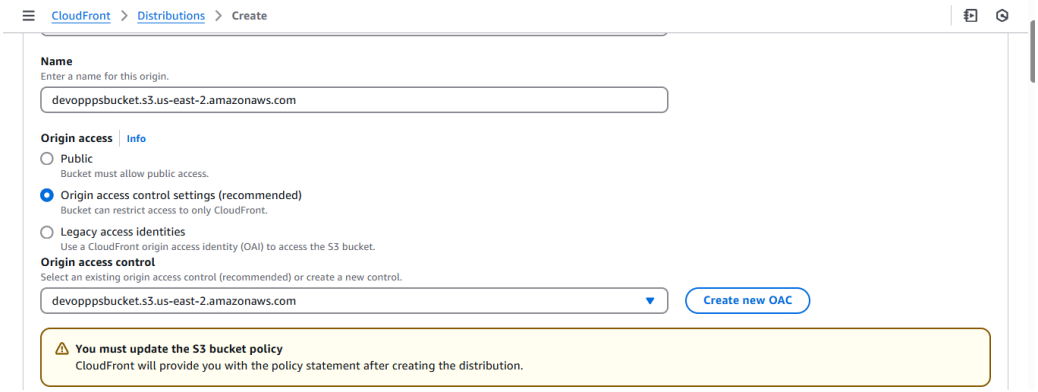

  - Set a new OAC using your S3website endpoint.
  - Click **Create new OAC** to create your OAC.

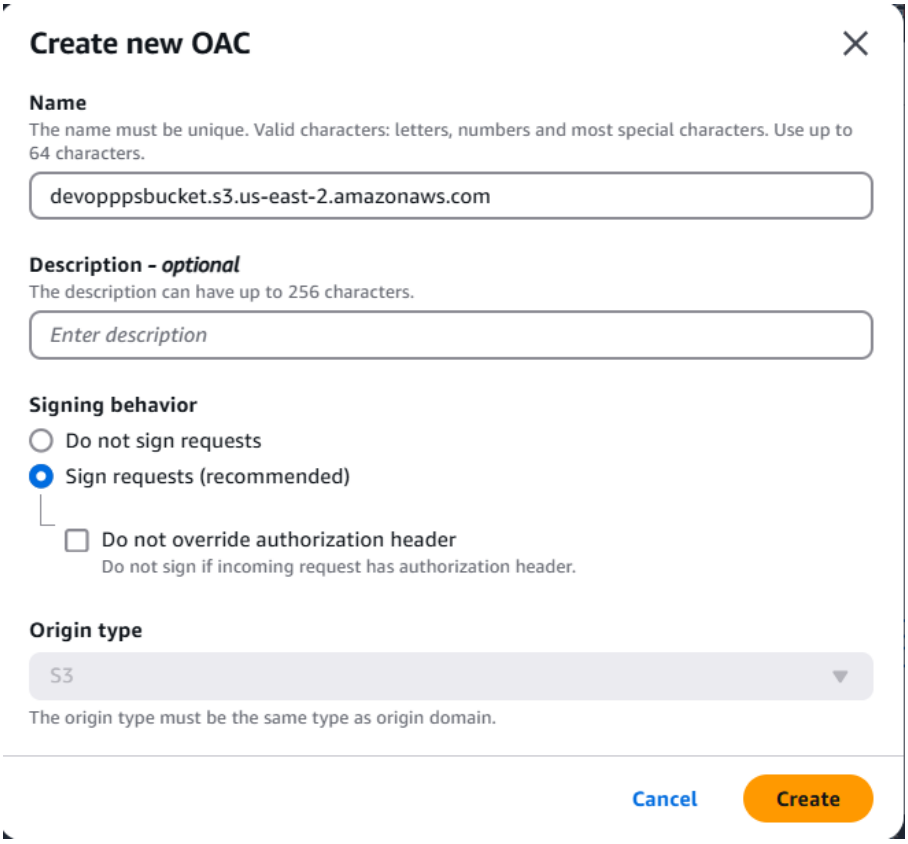

- Scroll down to default root objectand type in **index.html**
- Check "**Do not enable security protections**"
- Click on **Create Distribution**
- Update S3 Bucket policy by copying the policy to S3

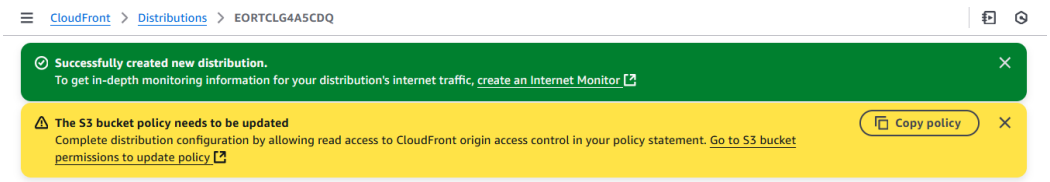

- Navigate to **Permissions**
- Check "**Block all public access**" to make the S3 bucket private.

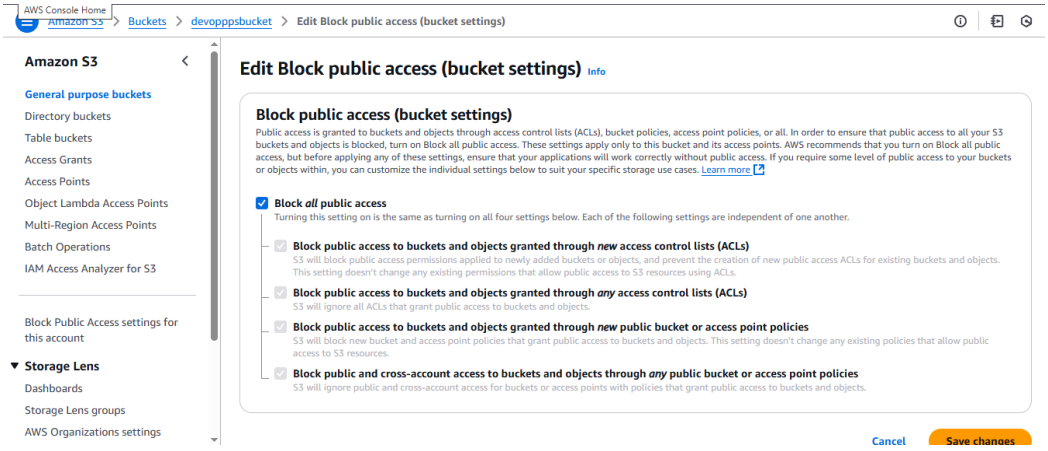

- Save changes
- Click on **Edit** policy
- Paste the policy you copied from CloudFront to access S3 Bucket.

Policy

```
{
    "Version": "2008-10-17",
    "Id": "PolicyForCloudFrontPrivateContent",
    "Statement": [
        {
            "Sid": "AllowCloudFrontServicePrincipal",
            "Effect": "Allow",
            "Principal": {
                "Service": "cloudfront.amazonaws.com"
            },
            "Action": "s3:GetObject",
            "Resource": "arn:aws:s3:::devopppsbucket/*",
            "Condition": {
                "StringEquals": {
                  "AWS:SourceArn": "arn:aws:cloudfront::831926586806:distribution/EORT
                }
            }
        }
    ]
}
```

- Save changes.
- Navigate to CloudFront and scroll down to details.
- Copy the Distribution Domain and paste on a new tab
- Refresh the page and you can access your website

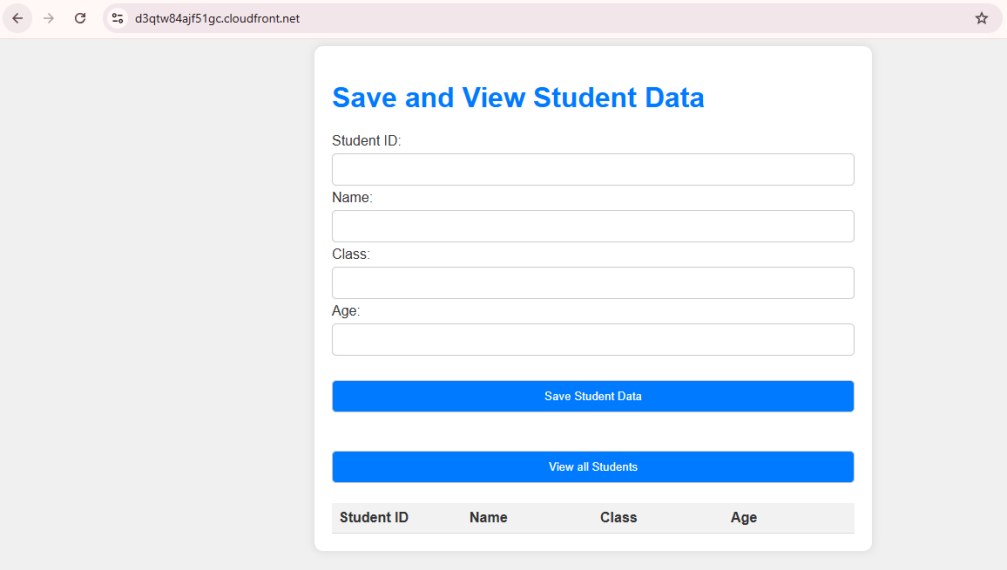

STEP 6: DELETE RESOURCES

- Delete API Gateway Resources
  - Go to Amazon API Gateway in the AWS Console
  - Select your API
  - Delete the API by clicking on Actions
  - Delete API
  - Confirm the deletion
- Delete AWS Lambda Function
  - Navigate to AWS Lambda
  - Select your functions used for the URL Shortener
  - Click Actions
  - Delete and Confirm
- Delete DynamoDB Table
  - Go to DynamoDB in the AWS Console
  - Find the table used for storing short URLs
  - Click Delete Table and Confirm
- Delete Amazon S3 Bucket
  - Open Amazon S3
  - Find the Bucket you created for storing the frontend
  - Empty the Bucket first by deleting the objects.
  - Then delete the Bucket

- Remove CloudFront Distribution
    - Go to Amazon CloudFront
    - Select the distribution associated with the URL Shortener
    - Disable it first and then Delete it
    - Wait for the status to change to Deleted